

Technical Implementation

Team Information

- **Team Member :** Pavan Sai Porapu
 - **Contact Email:** porapupavansai.22.csm@anits.edu.in
-

1. Overview of Technical Implementation

➤ What is the core functionality of your solution?

- The core functionality of this solution is to estimate **Air Quality Index (AQI)** based on various environmental and meteorological factors.
- The system uses past three days air quality index data based on factors such as concentration of pollutants (SO₂, O₃, NO₂, CO ect.) of metropolitan cities in india for training.
- The data is pre-processed using,
 - Numerical features - scaling (StandardScaling, MinMaxScaling)
 - Categorical features – OneHot Encoding
 - Feature expansion - Polynomial feature expansion of degree 2
- Different traditional and statistical machine learning models are trained to accurately predict the air quality index at a given location.

➤ How does your model integrate with the API data?

- **Training Phase:** The model is trained using historical data collected from weather APIs (Rapidapi.com)
- **Prediction Phase:** The system does not fetch real-time data but instead relies on user-inputted or pre-stored data to estimate AQI.

➤ What machine learning or deep learning techniques are used?

Trained multiple models such as both traditional ML and ensemble learning.

- **Regression Models:** Linear Regression, Ridge
- **Tree-Based Models:** Decision Tree, Random Forest, Gradient Boosting
- **Advanced Boosting Models:** XGBoost
- **Stacking Ensemble:** Random Forest, Decision Tree, and Gradient Boosting.
- **Hyperparameter Tuning:** GridSearchCV

Used Ensembled technique to get more accurate results. Stacking random forest regressor, Decision tree regressor and gradient boosting as final predictor.

2. ML Project Architecture

➤ What are the key stages in the end-to-end ML workflow?

(Describe the entire process from data collection to model deployment.)

End-to-End ML workflow :

1. Data Collection
 - Historical air pollution data are collected from weather APIs.
 - Data includes pollutant levels (PM2.5, PM10, NO2, CO, etc.) & location.
2. Data Preprocessing
 - Handle missing values by imputation or removal.
 - Convert timestamps to datetime format and extract relevant features (e.g., time of day, season).
 - Normalize data to a certain range to ensure consistency.
3. Feature Engineering
 - Generate polynomial features to capture non-linear relationships.
 - Encode categorical data if necessary.
 - Understanding Spread of data using Visualizations.
4. Model Training
 - Train various models (Linear Regression, Decision Trees, Random Forest, Gradient Boosting, XGBoost, etc.).
 - Use cross-validation (5-fold) to evaluate generalization performance.
 - Perform hyperparameter tuning using GridSearchCV.
5. Model Evaluation
 - Use metrics like Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and R² score.
 - Compare models and choose the best-performing one.
6. Model Deployment
 - Model is integrated in a web- interface and deployed (Streamlit)
7. Monitoring & Maintenance (Optional for Future Scalability)
 - Expanding model's knowledge by training various locations data.

➤ **How does data flow through the system?**

(Explain how raw data is collected, processed, stored, and used for training.)

1. Data Collection
 - Data is collected from an api endpoint and stored in .csv file format.
 - Data includes pollutants (PM2.5, PM10, NO2, etc.).
2. Data Processing & Storage
 - Handle missing values, convert timestamps, and normalize using StandardScaler.
 - Store cleaned data in CSV files for training.
3. Model Training
 - Data is split into training and testing sets.
 - Polynomial features are generated, and models are trained using regression and ensemble techniques.
4. Prediction & Output
 - The trained models is saved and later used with user-inputted data for AQI prediction.

➤ **What are the main components of the ML pipeline?**

(Break down data acquisition, preprocessing, feature engineering, model training, evaluation, and deployment.)

- Data Acquisition – Collect historical air pollutant data from APIs.
- Preprocessing – Clean data, handle missing values, normalize, and transform features.
- Feature Engineering – Generate polynomial features and select key variables.
- Model Training – Train multiple models (Random Forest, XGBoost, etc.) with hyperparameter tuning.
- Evaluation – Measure performance using RMSE, R^2 , and cross-validation.
- Deployment – Save the best model for making AQI predictions from user inputs.

3. Model Selection & Justification

- **Which model(s) did you choose and why? (e.g., Logistic Regression, SVM, Random Forest, Neural Networks, etc.)**
 1. Trained multiple regression models randomly and based on observations of the evaluation metrics and performance of the models ,decided to use an ensemble learning technique.
 2. We stacked three models Gradient Boosting, Decision Tree regressor, Random Forest regressor ,obtained an accurate model with R-squared score of 0.98 on both training and testing data.
- **Were multiple models tested? If yes, what comparisons were made?**

Performance of Models on Training Data:

	Model	Mean_absolute_error	Mean_squared_error	Root_Mean_Squared_error	R2_score
0	base_model(mean)	36.23	2,428.07	49.28	0
1	Linear Regression	10.37	315.43	17.76	0.87
2	DecisionTreeRegressor	11.41	312.88	17.69	0.87
3	Random Forest Regressor	1.26	61.57	7.85	0.97
4	Gradient Bosting	0.82	1.66	1.29	1
5	Ensembled_stacked_model	2.01	43.09	6.56	0.98

Figure 1:Evaluation metrics of models on training data

Performance of Models on Testing Data:

	Model	Mean_absolute_error	Mean_squared_error	Root_Mean_Squared_error	R2_score
0	base_model(mean)	36.23	2,428.07	49.28	0
1	Linear Regression	12.65	624.17	24.98	0.81
2	DecisionTreeRegressor	15.7	1,948.59	44.14	0.39
3	Random Forest Regressor	5.13	1,153.17	33.96	0.64
4	Gradient Bosting	3.45	330.03	18.17	0.9
5	Ensembled_stacked_model	3.45	68.56	8.28	0.98

Figure 2: Evaluation metrics of models on testing data

- **Was transfer learning or pre-trained models used?**

No Transfer learning or pre-trained models are applicable.

4. Training Process

- **How was the data split for training, validation, and testing?**
 - The data is divided into two splits
 1. Testing data -- 80%
 2. Training data – 20%
 - This is Achieved by using scikit-learn library using train test split function

- **What training techniques were used (e.g., data augmentation, early stopping)?**
 - Hyperparameter tuning – Gridsearchcv
 - Regularisation – to prevent overfitting to training data
 - Ensembled learning technique

- **What loss function and optimization algorithms were used (e.g., Adam, SGD)?**
 - Loss Function: Mean Squared Error (MSE) for regression models.
 - Optimization Algorithms:
 1. Gradient Descent variants (used in models like Ridge, Lasso, ElasticNet).
 2. Boosting algorithms (XGBoost, Gradient Boosting) use custom tree-based optimization.

- **Did you apply hyperparameter tuning? If so, what method was used (GridSearch, RandomSearch, Bayesian Optimization)?**
 - Yes, applied.
 - Method Used: GridSearchCV for exhaustive search over parameter combinations.
 - Models Tuned: Random Forest, Gradient Boosting (learning rate, depth, estimators).

5. Model Evaluation & Performance Metrics

- **What performance metrics were used for evaluation (e.g., Accuracy, Precision, Recall, F1 Score, RMSE)?**

Since this is a **regression problem**, the following metrics were used:

- Mean Absolute Error(MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- R² Score

➤ **How does your model compare to a baseline model?**

Performance of Models on Training Data:					
	Model	Mean_absolute_error	Mean_squared_error	Root_Mean_Squared_error	R2_score
0	base_model(mean)	36.23	2,428.07	49.28	0
1	Linear Regression	10.37	315.43	17.76	0.87
2	DecisionTreeRegressor	11.41	312.88	17.69	0.87
3	Random Forest Regressor	1.26	61.57	7.85	0.97
4	Gradient Bosting	0.82	1.66	1.29	1
5	Ensembled_stacked_model	2.01	43.09	6.56	0.98

Figure 3: Comparison of basedline model and other ml models

➤ **What were the key takeaways from the evaluation?**

1. Ensemble Models Performed Best
 - Stacked models (Random Forest + Gradient Boosting) outperformed individual models.
2. Feature Engineering Improved Accuracy
 - Polynomial features helped capture complex relationships, improving performance.
3. Hyperparameter Tuning was Essential
 - GridSearchCV optimized parameters, significantly reducing RMSE.
4. MAE & RMSE Showed Low Prediction Errors
 - The model made reasonably accurate predictions with minimal large deviations.

6. Optimization & Engineering Enhancements

Performance Optimization Techniques

- **Was model compression or pruning used?**
- Yes Used, techniques like tree pruning (for Decision Trees & Random Forests)
- **Did you implement quantization or reduced precision techniques?**
- Not implemented, but could improve inference speed.
- **How was inference speed optimized?**
- Used StandardScaler for fast data transformation.
 - Stacking Regressor optimized predictions by combining multiple models.

Code Efficiency & Engineering Best Practices

- **Was the model deployed in an optimized environment (e.g., GPU acceleration, parallel processing)?**

Model is going to be deployed in a free version in github or huggingface spaces.

- **What libraries or frameworks were used for performance improvements?**

- Scikit-learn, XGBoost for efficient ML operations.
- Pandas & NumPy for fast data handling.
- Streamlit for developing user- friendly interface.

- **How was code structured for maintainability and scalability?**

- Modular approach (separate data processing, training, and evaluation steps).

7. Model Deployment & Integration

- **Was the model deployed? (e.g., Flask API, FastAPI, Streamlit, Cloud services) **Provide the link to access your project****

- No the model is not deployed yet

- **How does the model interact with the user interface or API?**

- The model interact with the user through an interface developed using streamlit

- **What challenges were faced during deployment and how were they resolved?**

8. Challenges & Solutions

- **What were the biggest technical challenges faced during implementation?**

- Selecting the Best Model -- Balancing accuracy & inference speed was difficult.
- Hyperparameter Tuning Complexity -- Finding optimal settings for models like XGBoost took time.
- Time constraint – Due to lack of time I skipped some of the things which are may be essential.

- **How did your team overcome them?**

- Model Evaluation & Stacking – Combined models to improve accuracy.
- GridSearchCV for Hyperparameter Tuning – Optimized parameters efficiently.
- Prioritize the tasks at hand to complete due to time constraint.

9. Supporting Code & References (If applicable)

- **Attach or provide links to code snippets showcasing technical implementation.**
- **Mention any references or papers that supported your implementation.**